



CS263: Rust vs C++



Subramaniam Shankar
Rakshith Gopalakrishna

Introduction (Rust)

- Rust is developed by Mozilla.
- Rust is focused on providing performance and safety, especially safe concurrency and memory safety without garbage collection.
- Rust is very similar to C++, hence the comparison.
- Rust is the "most loved programming language" in the Stack Overflow developer survey since 2016.

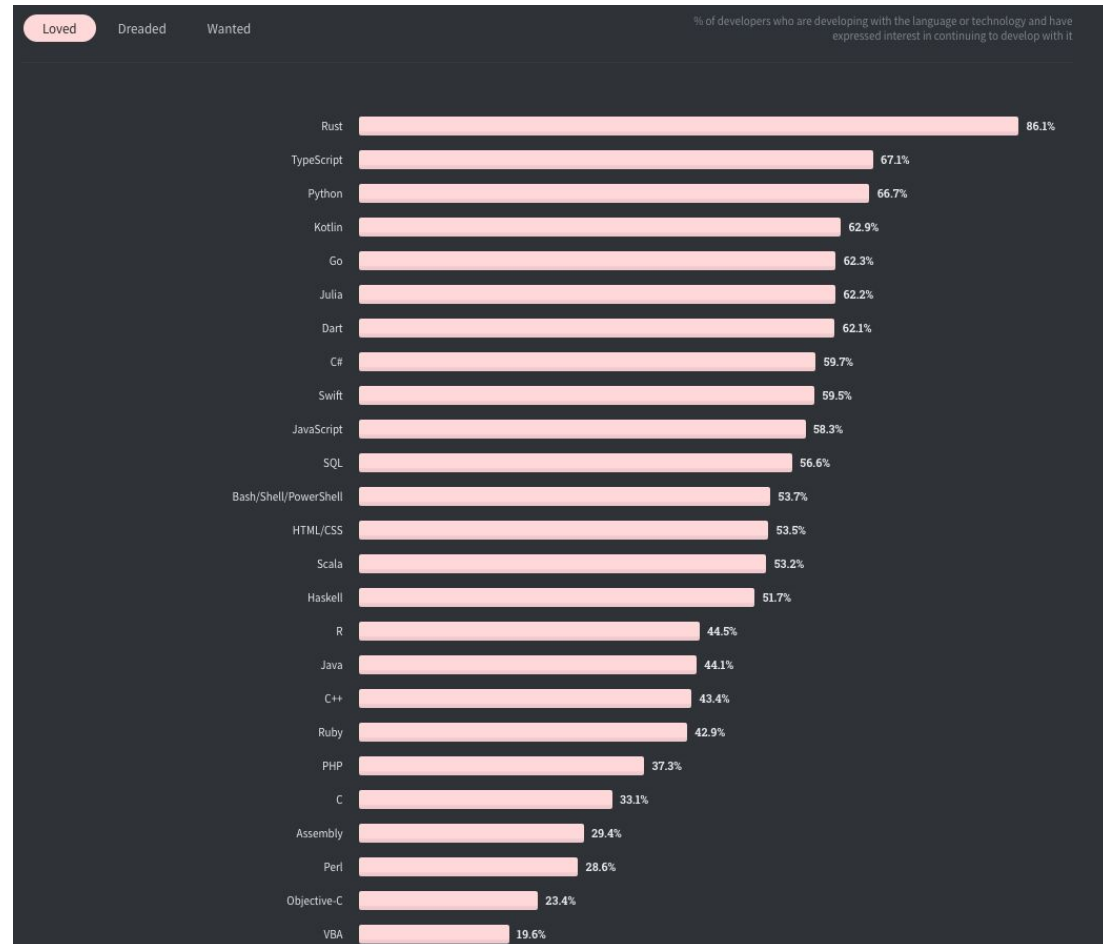
Introduction (Rust)

Rust: 86.1%

C++: 43.4%

Link:

<https://insights.stackoverflow.com/survey/2020#technology-most-loved-dreaded-and-wanted-languages-loved>



Ownership

- Each value in Rust has a variable that's called its owner.
- There can only be one owner at a time.
- When the owner goes out of scope, the value will be dropped.

```
{  
    let s = String::from("hello"); // s is valid from this point forward  
  
    // do stuff with s  
}  
// this scope is now over, and s is no  
// longer valid, call drop
```

Ownership (Move semantics)

Rust:

```
let s1 = String::from("hello"); // all resources to s1 allocated
                                // s1 valid from this point onward

let s2 = s1;
println!("{}", world, s1); // error
```


C++:

```
5 int main(){
6     std::string a = "Hello world";
7     std::cout<<a<<std::endl;
8     std::vector<std::string> vec;
9
10    vec.push_back(std::move(a));
11
12    std::cout<<vec[0]<<std::endl;
13    std::cout<<&vec[0]<<std::endl;
14    std::cout<<a<<std::endl;
15    a="New hello world";
16    std::cout<<&a<<std::endl;
17 }
```

```
0x7ffee651758
Hello world
0x7f811f404220
0x7ffee651758
```

Vector move vs. copy

A vector of 10,000,000 integers was moved/copied.



References in Rust

```
fn main() {  
    let s = String::from("hello"); // s comes into scope  
  
    takes_ownership(s);           // s's value moves into the function...  
                                  // ... and so is no longer valid here  
  
    let x = 5;                    // x comes into scope  
  
    makes_copy(x);                // x would move into the function,  
                                  // but i32 is Copy, so it's okay to still  
                                  // use x afterward  
  
} // Here, x goes out of scope, then s. But because s's value was moved, nothing  
  // special happens.  
  
fn takes_ownership(some_string: String) { // some_string comes into scope  
    println!("{}", some_string);  
} // Here, some_string goes out of scope and `drop` is called. The backing  
  // memory is freed.  
  
fn makes_copy(some_integer: i32) { // some_integer comes into scope  
    println!("{}", some_integer);  
} // Here, some_integer goes out of scope. Nothing special happens.
```

References in Rust (contd.)

```
fn main() {  
    let s1 = String::from("hello");  
  
    let len = calculate_length(&s1);  
  
    println!("The length of '{}' is {}.", s1, len);  
}  
  
fn calculate_length(s: &String) -> usize {  
    s.len()  
}
```


References in Rust (contd.)

```
fn main() {  
    let reference_to_nothing = dangle();  
}  
  
fn dangle() -> &String { // dangle returns a reference to a String  
  
    let s = String::from("hello"); // s is a new String  
  
    &s // we return a reference to the String, s  
} // Here, s goes out of scope, and is dropped. Its memory goes away.  
// Danger!
```

Invariants of Rust

- At any given time, you can have either one mutable reference or any number of immutable references.
- References must always be valid.

Smart pointers in C++: unique_ptr and Ownership

```
4  int main() {
5      std::unique_ptr<int> u;
6      int *p = new int;
7      *p=5;
8
9      u=std::unique_ptr<int>(p);
10
11     std::cout<<p<<" "<<*p<<std::endl;
12     std::cout<<u.get()<<" "<<*u<<std::endl;
13
14     std::unique_ptr<int> u2;
15     u2 = std::unique_ptr<int>(std::move(u));
16
17     std::cout<<u.get()<<std::endl;
18     std::cout<<u2.get()<<" "<<*u2<<std::endl;
19
20     std::unique_ptr<int> u3(p);
21     std::cout<<u3.get()<<" "<<*u3<<std::endl;
22 }
```

```
0x562d369dfcb0 5
0x562d369dfcb0 5
0
0x562d369dfcb0 5
0x562d369dfcb0 5
free(): double free detected in tcache 2
[1] 976339 abort (core dumped) bin/unique_ptr
```

Smart pointer in Rust: Box pointer

```
{ 2  fn main() {  
    3      let a = Box::new(5);  
    4      println!("result: {:?}", *a);  
} 5  }
```

Box pointer from Raw pointers (unsafe Rust)

```
{ 1 fn main() {  
2     let a = Box::into_raw(Box::new(String::from("Hello World")));  
3  
4     let u = Box::new(a);  
5     println!("{:p} {:?}", a, unsafe{&*a});  
6     println!("{:p} {:?} {:?}", u, *u, unsafe{&*(u)});  
7  
8     let u4 = unsafe{Box::from_raw(a)};  
9     println!("{:?}", u4);  
10  
11     let u5 = unsafe{Box::from_raw(a)};  
12     println!("{:?}", u5);  
13  
14     println!("{:p} {:?}", a, unsafe{&*a});  
15  
16     // unsafe{  
17     //     drop(Box::from_raw(a));  
18     // };  
19 }
```

0x7fa9d9c05c10 "Hello World"

0x7fa9d9c05c30 0x7fa9d9c05c10 "Hello World"

"Hello World"

"Hello World"

0x7fa9d9c05c10 "Hello World"


box(13199,0x109583dc0) malloc: *** error for object 0x7fa9d9c05c00: pointer being freed was not allocated

box(13199,0x109583dc0) malloc: *** set a breakpoint in malloc_error_break to debug

Abort trap: 6

Linked List (Single)

Variants of Single Linked List



Smart pointers in C++: shared_ptr

```
4  int main(){
5      std::shared_ptr<int> u;
6      int *p = new int;
7      *p=5;
8
9      u=std::shared_ptr<int>(p);
10
11     std::cout<<p<<" "<<*p<<std::endl;
12     std::cout<<u.get()<<" "<<*u<<std::endl;
13
14     std::shared_ptr<int> u2;
15     u2 = std::shared_ptr<int>(std::move(u));
16
17     std::cout<<u.get()<<std::endl;
18     std::cout<<u2.get()<<" "<<*u2<<std::endl;
19
20     std::shared_ptr<int> u3(u2);
21     std::cout<<u3.get()<<" "<<*u3<<std::endl;
22
23     std::shared_ptr<int> u4(p);
24     std::cout<<u4.get()<<" "<<*u4<<std::endl;
25
26 }
```

```
0x5567f139aeb0 5
0x5567f139aeb0 5
0
0x5567f139aeb0 5
0x5567f139aeb0 5
0x5567f139aeb0 5
free(): double free detected in tcache 2
[1] 1068980 abort (core dumped) bin/shared_ptr
```

Memory leak with shared_ptr

```
int main(){
    Node *a = new Node(1);
    Node *b = new Node(2);
    Node *c = new Node(3);
    Node *d = new Node(4);

    std::shared_ptr<Node> a_ptr, b_ptr, c_ptr, d_ptr;
    a_ptr = std::make_shared<Node>(*a);
    b_ptr = std::make_shared<Node>(*b);
    c_ptr = std::make_shared<Node>(*c);
    d_ptr = std::make_shared<Node>(*d);

    a->next = b_ptr;
    b->prev = a_ptr;

    b->next = c_ptr;
    c->prev = b_ptr;

    c->next = d_ptr;
    d->prev = c_ptr;
```

```
==1155432== Memcheck, a memory error detector
==1155432== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==1155432== Using Valgrind-3.16.0.GIT and LibVEX; rerun with -h for copyright info
==1155432== Command: bin/shared_ptr_mem_leak
==1155432==
==1155432==
==1155432== HEAP SUMMARY:
==1155432==    in use at exit: 384 bytes in 8 blocks
==1155432==    total heap usage: 9 allocs, 1 frees, 73,088 bytes allocated
==1155432==
==1155432== LEAK SUMMARY:
==1155432==    definitely lost: 160 bytes in 4 blocks
==1155432==    indirectly lost: 224 bytes in 4 blocks
==1155432==    possibly lost: 0 bytes in 0 blocks
==1155432==    still reachable: 0 bytes in 0 blocks
==1155432==    suppressed: 0 bytes in 0 blocks
==1155432== Rerun with --leak-check=full to see details of leaked memory
==1155432==
==1155432== For lists of detected and suppressed errors, rerun with: -s
==1155432== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```


Rc (Reference Counted) Smart Pointer: Rust

```
11 fn main() {
12     let a = Rc::new(RefCell::new(Node{
13         value: 1,
14         prev: None,
15         next: None
16     }));
17
18     let b = Rc::new(RefCell::new(Node{
19         value: 2,
20         prev: None,
21         next: None
22     }));
23
24     let c = Rc::new(RefCell::new(Node{
25         value: 3,
26         prev: None,
27         next: None
28     }));
29
30     let d = Rc::new(RefCell::new(Node{
31         value: 4,
32         prev: None,
33         next: None
34     }));
35
36     (*a).borrow_mut().next = Some(Rc::clone(&b));
37
38     (*b).borrow_mut().prev = Some(Rc::clone(&a));
39     (*b).borrow_mut().next = Some(Rc::clone(&c));
40
41     (*c).borrow_mut().prev = Some(Rc::clone(&b));
42     (*c).borrow_mut().next = Some(Rc::clone(&d));
43
44     (*d).borrow_mut().prev = Some(Rc::clone(&c));
45
46     // println!("{:?}", a);
47 }
```

```
==16029== Memcheck, a memory error detector
==16029== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==16029== Using Valgrind-3.16.0.GIT and LibVEX; rerun with -h for copyright info
==16029== Command: ../bin/rc_mem_leak
==16029==
--16029-- run: /usr/bin/dsymutil "../bin/rc_mem_leak"
==16029==
==16029== HEAP SUMMARY:
==16029==    in use at exit: 14,632 bytes in 164 blocks
==16029==    total heap usage: 186 allocs, 22 frees, 19,641 bytes allocated
==16029==
==16029== LEAK SUMMARY:
==16029==    definitely lost: 48 bytes in 1 blocks
==16029==    indirectly lost: 144 bytes in 3 blocks
==16029==    possibly lost: 4,392 bytes in 5 blocks
==16029==    still reachable: 10,048 bytes in 155 blocks
==16029==    suppressed: 0 bytes in 0 blocks
==16029== Rerun with --leak-check=full to see details of leaked memory
==16029==
==16029== For lists of detected and suppressed errors, rerun with: -s
==16029== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 1 from 1)
```

weak_ptr(C++)/Weak<T>(Rust) to the rescue

```
class Node{
public:
    int value;
    // Node *next;
    std::shared_ptr<Node> next;
    std::weak_ptr<Node> prev;

    Node(int val){
        this->value = val;
    }
};
```

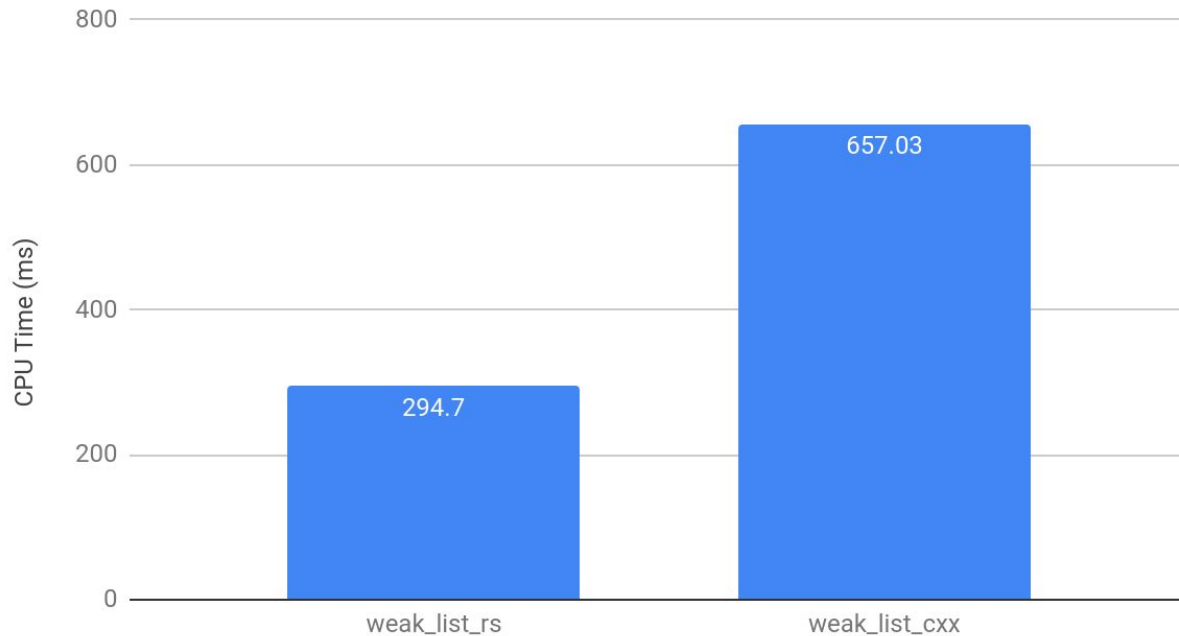
```
==1161301== Memcheck, a memory error detector
==1161301== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==1161301== Using Valgrind-3.16.0.GIT and LibVEX; rerun with -h for copyright info
==1161301== Command: bin/weak_list
==1161301==
10
==1161301==
==1161301== HEAP SUMMARY:
==1161301==     in use at exit: 0 bytes in 0 blocks
==1161301==   total heap usage: 10,000,004 allocs, 10,000,004 frees, 320,073,792 bytes allocated
==1161301==
==1161301== All heap blocks were freed -- no leaks are possible
==1161301==
==1161301== For lists of detected and suppressed errors, rerun with: -s
==1161301== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

```
1 use std::cell::RefCell;
2 use std::rc::{Rc, Weak};
3
4 #[derive(Debug)]
5 struct Node {
6     value: i32,
7     next: Option<Rc<RefCell<Node>>>,
8     prev: Option<Weak<RefCell<Node>>>
9 }
10
11 #[derive(Debug)]
12 struct List {
13     head: Option<Rc<RefCell<Node>>>
14 }
```

```
==1169227== Memcheck, a memory error detector
==1169227== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==1169227== Using Valgrind-3.16.0.GIT and LibVEX; rerun with -h for copyright info
==1169227== Command: bin/weak_list
==1169227==
10
==1169227==
==1169227== HEAP SUMMARY:
==1169227==     in use at exit: 0 bytes in 0 blocks
==1169227==   total heap usage: 5,000,019 allocs, 5,000,019 frees, 240,003,329 bytes allocated
==1169227==
==1169227== All heap blocks were freed -- no leaks are possible
==1169227==
==1169227== For lists of detected and suppressed errors, rerun with: -s
==1169227== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

Weak_ptr to the rescue

Doubly linked list with weak pointers



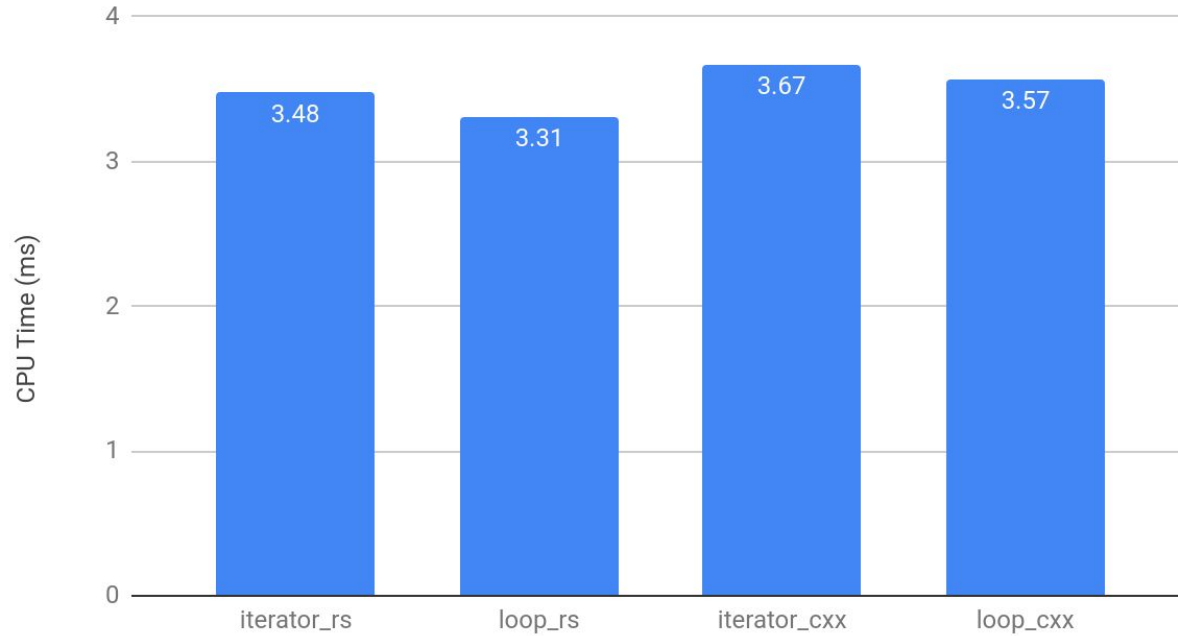
Zero cost abstractions (iterators and closures)

```
{ 7  fn main() {  
8      let mut a = Vec::new();  
9      let mut b = Vec::new();  
10  
11      for i in 1..=50000 {  
12          a.push(i);  
13          b.push(10000 - i);  
14      }  
15  
16      let numbers = a.iter()  
17          .zip(b.iter())  
18          .map(|(a, b)| a * b)  
19          .filter(|a| *a > 5000)  
20          .take(4)  
21          .collect::<Vec<_>>();  
22  
23      println!("{:?}", numbers);  
24  }  
25 }
```

```
{ 8  fn main() {  
9  
10      let mut a = Vec::new();  
11      let mut b = Vec::new();  
12  
13      for i in 1..=50000 {  
14          a.push(i);  
15          b.push(10000 - i);  
16      }  
17  
18      let mut numbers = Vec::new();  
19  
20      for i in 0..min(a.len(), b.len()) {  
21          let product = a[i] * b[i];  
22  
23          if product > 5000 {  
24              numbers.push(product);  
25          }  
26  
27          if numbers.len() == 4 {  
28              break;  
29          }  
30      }  
31  
32      println!("{:?}", numbers);  
33  }
```

Zero cost abstractions

Zero cost abstractions



Concurrency: Sharing Same memory

```
int main(){
    std::shared_ptr<int> ptr = std::make_shared<int>(2011);

    std::vector<std::thread*> thread_vec;
    for (int i= 0; i<10; i++){
        std::thread *t = new std::thread([]()mutable{
            std::shared_ptr<int> local(ptr);
            *local = i;
            std::cout<<"Thread "<<i<<" "<<local.get()<<" "<<*local\
            <<" "<<"Original ptr "<<ptr.get()<<" "<<*ptr<<std::endl;
        });
        thread_vec.push_back(t);
    }
    for(auto i:thread_vec){
        i->join();
    }
    std::cout<<"Original pointer"<<std::endl;
    std::cout<<ptr.get()<<" "<<*ptr<<std::endl;
}
```

```
int main(){
    std::shared_ptr<int> ptr = std::make_shared<int>(2011);


    std::vector<std::thread*> thread_vec;
    for (int i= 0; i<10; i++){
        std::thread *t = new std::thread([&]()mutable{
            std::shared_ptr<int> local(ptr);
            *local = i;
            std::cout<<"Thread "<<i<<" "<<local.get()<<" "<<*local\
            <<" "<<"Original ptr "<<ptr.get()<<" "<<*ptr<<std::endl;
        });
        thread_vec.push_back(t);
    }
    for(auto i:thread_vec){
        i->join();
    }
    std::cout<<"Original pointer"<<std::endl;
    std::cout<<ptr.get()<<" "<<*ptr<<std::endl;
}
```


Concurrency: Sharing Same memory

```
fn main() {  
    let ptr = Arc::new(Mutex::new(2011));  
    let mut thread_vec = vec![];  
  
    for i in 0..10 {  
        let ptr = Arc::clone(&ptr);  
        let handle = thread::spawn (move || {  
            let mut val = (*ptr).lock().unwrap();  
            *val = i;  
            println!("Thread {:?} {:p} {:?} Original ptr {:p} {:?}", i, ptr, *val, ptr, *val);  
        });  
        thread_vec.push(handle);  
    }  
  
    for thread in thread_vec {  
        thread.join().unwrap();  
    }  
  
    println!("Original pointer\n{:p} {:?}", ptr, (*ptr).lock().unwrap());  
}
```

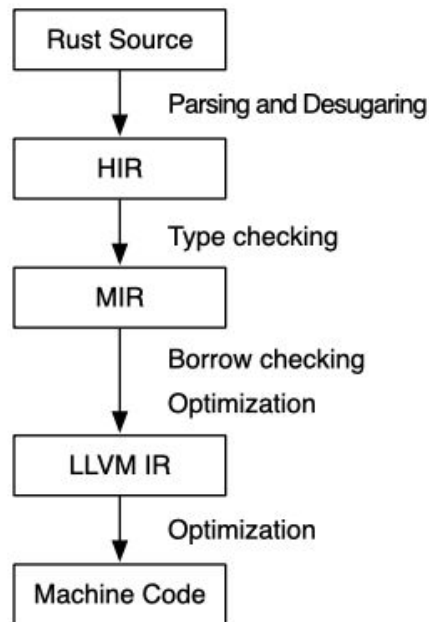
Concurrency: Message Passing

MPSC Time



Other Features that we explored

- Monomorphization of Generic types in both C++ and Rust by examining the assembly
- Dynamic Dispatch
- Index Loop unrolling



Questions?